

Metaprogramming Applied to Numerical Problems

A Generic Implementation of Runge-Kutta Algorithms

Mario Mulansky
Karsten Ahnert

University of Potsdam

April 16, 2012



Content

- Ordinary Differential Equations (ODEs)
- numerical routines: Runge-Kutta Schemes
- “classic” implementations
- generic implementation using Template Metaprogramming
- performance tests
- summary

Ordinary Differential Equations

ODEs are the typical way to describe physical, biological, chemical, ... processes and thus play a fundamental role in mathematical modelling.

- Newton's equation of motion
- Reaction-diffusion systems
- Modelling of interacting neuronal networks

Ordinary Differential Equations

ODEs are the typical way to describe physical, biological, chemical, ... processes and thus play a fundamental role in mathematical modelling.

- Newton's equation of motion
- Reaction-diffusion systems
- Modelling of interacting neuronal networks

Also, ODEs are used as approximations to Partial Differential Equations for numerical treatments.

Ordinary Differential Equations

A first order ODE is written in its most general form as:

$$\frac{d}{dt}\vec{x}(t) = \vec{f}(\vec{x}, t) \quad (1)$$

- $\vec{x}(t)$ is the function in demand (here: trajectory)
- t is the independent variable (here: time)
- $f(x, t)$ is the rhs, governing the behavior of x

Initial Value Problem (IVP):

$$\dot{x} = f(x, t), \quad x(t = 0) = x_0 \quad (2)$$

Examples

- $\dot{x} = -\lambda x$ solution: $x(t) = x_0 e^{-\lambda t}$
- $\ddot{x} = \omega^2 x \rightarrow \begin{cases} \dot{x} = p \\ \dot{p} = -\omega^2 x \end{cases}$ solution: $x(t) = A \sin(\omega t + \varphi_0)$.

Examples

- $\dot{x} = -\lambda x$ solution: $x(t) = x_0 e^{-\lambda t}$
- $\ddot{x} = \omega^2 x \rightarrow \begin{cases} \dot{x} = p \\ \dot{p} = -\omega^2 x \end{cases}$ solution: $x(t) = A \sin(\omega t + \varphi_0)$.
 $\dot{x} = \sigma(y - x)$
- Lorenz System: $\dot{y} = x(R - z) - y$ solution: ?
 $\dot{z} = xy - \beta z$.

Chaotic system (for certain parameter values σ, R, β), hence the solution can not be written in analytic form.

Examples

- $\dot{x} = -\lambda x$ solution: $x(t) = x_0 e^{-\lambda t}$
- $\ddot{x} = \omega^2 x \rightarrow \begin{cases} \dot{x} = p \\ \dot{p} = -\omega^2 x \end{cases}$ solution: $x(t) = A \sin(\omega t + \varphi_0)$.
- $\dot{x} = \sigma(y - x)$
- Lorenz System: $\dot{y} = x(R - z) - y$ solution: ?
 $\dot{z} = xy - \beta z$.

Chaotic system (for certain parameter values σ, R, β), hence the solution can not be written in analytic form.

\Rightarrow numerical methods to solve ODEs are required for more complicated systems.

Runge-Kutta Scheme


One class of algorithms to solve IVP of ODEs.

- Discretized time $t \rightarrow t_n = t_0 + n \cdot h$ with (small) time step h
- Trajectory $x(t) \rightarrow x_n \approx x(t_n)$
- Iteration along trajectory: $x_0 \longrightarrow x_1 \longrightarrow x_2 \dots$
- One-step method: $x_1 = \Phi(x_0)$, $x_2 = \Phi(x_1)$, \dots

Runge-Kutta Scheme

One class of algorithms to solve IVP of ODEs.

- Discretized time $t \rightarrow t_n = t_0 + n \cdot h$ with (small) time step h
- Trajectory $x(t) \rightarrow x_n \approx x(t_n)$
- Iteration along trajectory: $x_0 \rightarrow x_1 \rightarrow x_2 \dots$
- One-step method: $x_1 = \Phi(x_0)$, $x_2 = \Phi(x_1)$, ...



Runge-Kutta Scheme

Numerically solve the Initial Value Problem (IVP) of the ODE:

$$\dot{x}(t) = f(x, t), \quad x(t = 0) = x_0. \quad (3)$$

A Runge-Kutta scheme with s stages and parameters $c_1 \dots c_s$, $a_{21}, a_{31}, a_{32}, \dots, a_{ss-1}$ and $b_1 \dots b_s$ gives the approximate solution for $x_1 \approx x(h)$ starting at x_0 by computing:

$$x_1 = x_0 + h \sum_{i=1}^s b_i F_i \quad \text{where} \quad F_i = f\left(x_0 + h \sum_{j=1}^{i-1} a_{ij} F_j, h c_i\right). \quad (4)$$

This approximate solution x_1 is exact up to some order p . Repeating the whole procedure brings you from x_1 to x_2 , then to x_3 and so on.

At each stage i the following calculations have to be performed ($y_1 = x_0$) :

$$F_i = f(y_i, hc_i), \quad y_{i+1} = x_0 + h \sum_{j=1}^i a_{i+1,j} F_j, \quad i = 1 \dots s-1$$
$$F_s = f(y_s, hc_s), \quad x_1 = x_0 + h \sum_{j=1}^s b_j F_j.$$

The parameters a , b and c define the so-called Butcher tableau.

Butcher Tableau

Parameters a , b , and c are typically written as Butcher tableau:

c_1					
c_2	$a_{2,1}$				
c_3	$a_{3,1}$	$a_{3,2}$			
\vdots	\vdots		\ddots		
c_s	$a_{s,1}$	$a_{s,2}$	\dots	$c_{s,s-1}$	
	b_1	b_2	\dots	b_{s-1}	b_s

The Butcher Tableau fully defines the Runge-Kutta scheme. Each line of the tableau represents one stage of the scheme.

Explicit Non-Generic Implementation

Given parameters c_i , a_{ij} , b_i

```
F_1 = f( x , t + c_1*dt );  
x_tmp = x + dt*a_21 * F_1;  
  
F_2 = f( x_tmp , t + c_2*dt );  
x_tmp = x + dt*a_31 * F_1 + dt*a_32 * F_2;  
  
// ...  
  
F_s = f( x_tmp , t + c_s*dt );  
x_end = x + dt*b_1 * F_1 + dt*b_2 * F_2 + ...  
          + dt*b_s * F_s;
```

Not generic: Each stage written hard coded – you have to adjust the algorithm when implementing a new scheme.

Run Time Implementation

Given parameters $a[] []$, $b[]$, $c[]$.

```
F[0] = f( x , t + c[0]*dt );
x_tmp = x + dt*a[0][0] * F[0];

for( int i=1 ; i<s-1 ; ++i )
{
    F[i] = f( x_tmp , t + c[i]*dt );
    x_tmp = x;
    for( int j=0 ; j<i+1 : ++j )
        x_tmp += dt*a[i][j] * F[j];
}

F[s-1] = f( x_tmp , t + c[s-1]*dt );
x_end = x;
for( int j=0 ; j<s : ++j )
    x_end += dt*b[j] * F[j];
```

Run Time Implementation

Given parameters $a[] []$, $b[]$, $c[]$.

```
F[0] = f( x , t + c[0]*dt );
x_tmp = x + dt*a[0][0] * F[0];

for( int i=1 ; i<s-1 ; ++i )
{
    F[i] = f( x_tmp , t + c[i]*dt );
    x_tmp = x;
    for( int j=0 ; j<i+1 : ++j )
        x_tmp += dt*a[i][j] * F[j];
}

F[s-1] = f( x_tmp , t + c[s-1]*dt );
x_end = x;
for( int j=0 ; j<s : ++j )
    x_end += dt*b[j] * F[j];
```

Generic, but factor 2 slower than explicit implementation!

Why Bad Performance

The run time generic code is hard to optimize for the compiler, because:

- Double `for` loop with inner bound depending on outer loop variable.
- 2D array `double** a` must be dynamically allocated:

```
a = new double*[s];  
for( int i=0 ; i<s ; ++i )  
    a[i] = new double[i+1];  
a[0][0] = ...;  
a[1][0] = ...; a[1][1] = ...;  
...
```

→ lives on heap, harder to be optimized compared to stack.

- Many more issues possible (optimizers are rather complex).

What to do?

Idea:

Use template engine to generate code that can be efficiently optimized by the Compiler.

What to do?

Idea:

Use template engine to generate code that can be efficiently optimized by the Compiler.

More specifically, we will use Template Metaprogramming to:

- Generate fixed size arrays: `a_1[1]` , `a_2[2]` , ... , `a_s[s]`
- Unroll the outer `for`-loop (over stages `s`) so the compiler sees sequential code.

As result, the code seen by the compiler/optimizer (after resolving templates) is very close to the non-generic version and thus as fast, hopefully.

Generic Runge-Kutta Algorithm

Idea:

- Write a Metaprogram that creates Runge-Kutta algorithms
- Metaprogram input: Parameters of the RK scheme (Butcher Tableau)
- Main objective: **Resulting program should be as fast as direct implementation**

With such a Metaprogram you can implement any new Runge-Kutta scheme by just providing the Butcher tableau.

- Decrease in programming time
- Less bugs
- Better maintainability

The Generic Implementation

Define a structure representing one stage of the Runge-Kutta scheme:

```
template< int i >
struct stage // general (intermediate) stage,  $i > 0$ 
{
    double c; // parameter  $c_i$ 
    array<double,i> a; // parameters  $a_{i+1,i} \dots a_{i,i}$ 
                        //  $b_1 \dots b_j$  for the last stage
};
```

Given an instance of this stage with c and a set appropriately the corresponding Runge-Kutta stage can be calculated.

The Generic Implementation

```
// x , x_tmp , t , dt and F defined outside
template< int i >
void calc_stage( const stage< i > &stage )
{ // performs the calculation of the i-th stage
  if( i == 1 ) // first stage?
    F[i-1] = f( x , t + stage.c * dt );
  else
    F[i-1] = f( x_tmp , t + stage.c * dt );

  if( i < s ) { // intermediate stage?
    x_tmp = x;
    for( int j=0 ; j<i : ++j )
      x_tmp += dt*stage.a[j] * F[j];
  } else { // last stage
    x_end = x;
    for( int j=0 ; j<i : ++j )
      x_end += dt*stage.a[j] * F[j];
  }
}
```

The Generic Implementation

Generate list of stage types: stage<1> , stage<2>, ... , stage<s>
using Boost.MPL (MetaProgramming Library) and Boost.Fusion.

```
typedef mpl::range_c< int , 1 , s > stage_indices;

typedef typename fusion::result_of::as_vector
< typename mpl::push_back
  < typename mpl::copy
    < stage_indices,
      mpl::inserter
    <
      mpl::vector0<> ,
      mpl::push_back< mpl::_1 , stage_wrapper< mpl::_2 , stage > >
    >
  >::type , stage< double , stage_count , last_stage >
>::type
>::type stage_vector_base; //fusion::vector< stage<1> , stage<2> , ... , stage<s>

struct stage_vector : stage_vector_base
{
  // initializer methods
  stage_vector( const a_type &a , const b_type &b , const c_type &c )
  {
    // ...
  }
}
```

The Generic Implementation

Parameter types for a, b and c:

```
typedef typename fusion::result_of::as_vector
< typename mpl::copy
  < stage_indices ,
    mpl::inserter
      < mpl::vector0< > ,
        mpl::push_back< mpl::_1 ,
          array_wrapper< double , mpl::_2 > >
      >
    >::type
>::type a_type; //fusion::vector< array<double,1> , array<double,2> , ... >

typedef array< double , s > b_type;
typedef array< double , s > c_type;
```


The Generic Implementation

Parameter types for a, b and c:

```
typedef typename fusion::result_of::as_vector
< typename mpl::copy
  < stage_indices ,
    mpl::inserter
      < mpl::vector0< > ,
        mpl::push_back< mpl::_1 ,
          array_wrapper< double , mpl::_2 > >
      >
    >::type
>::type a_type; //fusion::vector< array<double,1> , array<double,2> , ... >

typedef array< double , s > b_type;
typedef array< double , s > c_type;
```

Instead of a dynamically allocated `double**` the compiler/optimizer sees fixed size arrays: `array<double,1>` , `array<double,2>`, ...

→ **better optimization possibilities**

The Generic Implementation

The actual Runge-Kutta step (details omitted):

```
fusion::for_each( stages ,  
                  calc_stage_caller( f , x , x_tmp , x_end , F , t , dt ) );
```

Remember: `stages` is `fusion::vector< stage<1> , stage<2> , ... >`

For each of the `stages`, `calc_stage` gets called, but the `for_each`-loop is **executed by the compiler!**

The Generic Implementation

The actual Runge-Kutta step (details omitted):

```
fusion::for_each( stages ,  
                  calc_stage_caller( f , x , x_tmp , x_end , F , t , dt ) );
```

Remember: `stages` is `fusion::vector< stage<1> , stage<2> , ... >`

For each of the `stages`, `calc_stage` gets called, but the `for_each`-loop is **executed by the compiler!**

The compiler/optimizer sees sequential code:

```
calc_stage( stage_1 ); // stage_1 is an  
calc_stage( stage_2 ); // instance of stage<1>  
... // similar for stage_2 ...  
calc_stage( stage_s );
```

→ **better optimization possibilities**

The Generic Stepper

Provide some handy interface to the generic algorithm:

```
template< int s >
class generic_runge_kutta
{
public:
    generic_runge_kutta( const coef_a_type &a ,
                        const coef_b_type &b ,
                        const coef_c_type &c )
        : m_stages( a , b , c )
    { }

    void do_step( System f , const state_type &x , const double t ,
                 state_type &x_out , const double dt )
    {
        fusion::for_each( m_stages , calc_stage_caller( f , x , m_x_tmp , x_out ,
                                                         m_F , t , dt ) );
    }

private:
    stage_vector m_stages;
    state_type m_x_tmp;

protected:
    state_type m_F[s];
};
```

Example: Runge-Kutta 4

Butcher Tableau:

0				
0.5	0.5			
0.5	0	0.5		
1.0	0	0	1.0	
	1/6	1/3	1/3	1/6

```
// define the butcher array
const array< double , 1 > a1 = {{ 0.5 }};
const array< double , 2 > a2 = {{ 0.0 , 0.5 }};
const array< double , 3 > a3 = {{ 0.0 , 0.0 , 1.0 }};

const a_type a = fusion::make_vector( a1 , a2 , a3 );
const b_type b = {{ 1.0/6.0 , 1.0/3.0 , 1.0/3.0 , 1.0/6.0 }};
const c_type c = {{ 0.0 , 0.5 , 0.5 , 1.0 }};


// create the stages with the rk4 parameters a,b,c
generic_runge_kutta< 4 > rk4( a , b , c );
// do one rk4 step
rk4.do_step( lorenz , x , 0.0 , x , 0.1 );
```

Performance

Did we achieve our aim? Test RK4 on Lorenz System!

Performance

Did we achieve our aim? Test RK4 on Lorenz System!

**Processors:**

Intel Core i7 830

Intel Core i7 930

Intel Xeon X5650

Intel Core2Quad Q9550

AMD Opteron 2224

AMD PhenomII X4 945

Compilers:


gcc 4.3 , 4.4 , 4.5 , 4.6

intel icc 11.1 , 12.0

msvc 9.0

Performance

Did we achieve our aim? Test RK4 on Lorenz System!



Yes!

Processors:

Intel Core i7 830

Intel Core i7 930

Intel Xeon X5650

Intel Core2Quad Q9550

AMD Opteron 2224

AMD PhenomII X4 945

Compilers:


gcc 4.3 , 4.4 , 4.5 , 4.6

intel icc 11.1 , 12.0

msvc 9.0

Performance

Did we achieve our aim? Test RK4 on Lorenz System!



Processors:

Intel Core i7 830

Intel Core i7 930

Intel Xeon X5650

Intel Core2Quad Q9550

AMD Opteron 2224

AMD PhenomII X4 945

Compilers:

gcc 4.3 , 4.4 , 4.5 , 4.6

intel icc 11.1 , 12.0


msvc 9.0

Yes!

- On modern compilers (Intel 12, gcc 4.5/4.6) as fast as explicit code.
- Older compilers might produce slightly worse performant code.
- Always factor 2 better than run time generic implementation.

Performance

Second test with a different scheme: Runge-Kutta Cash-Karp 5(4)



Conclusions

We implemented a generic Runge-Kutta algorithm that executes **any** RK scheme and has the following properties:

- Parameters (Butcher Tableau) can be defined in a natural way as C++ Arrays
- By virtue of Template Metaprogramming our code is as fast as direct implementation of the specific scheme
- **Major improvement (factor 2) compared to generic run time implementation** (but some increase in compile time)
- Embedded methods with error estimate can also be easily covered in a generic way
- This technique can be applied to other numerical problems, e.g. spline fitting, ...

Conclusions

We implemented a generic Runge-Kutta algorithm that executes **any** RK scheme and has the following properties:

- Parameters (Butcher Tableau) can be defined in a natural way as C++ Arrays
- By virtue of Template Metaprogramming our code is as fast as direct implementation of the specific scheme
- **Major improvement (factor 2) compared to generic run time implementation** (but some increase in compile time)
- Embedded methods with error estimate can also be easily covered in a generic way
- This technique can be applied to other numerical problems, e.g. spline fitting, ...

Thank you

www.odeint.com